

BugRedux: Reproducing Field Failures for In-house Debugging

Wei Jin and Alessandro Orso

Georgia Institute of Technology

Email: {weijin|orso}@gatech.edu



Ming Xiao Yuan

Masters Student at Polytechnique Montreal



The problematic

I can't reproduce the error!



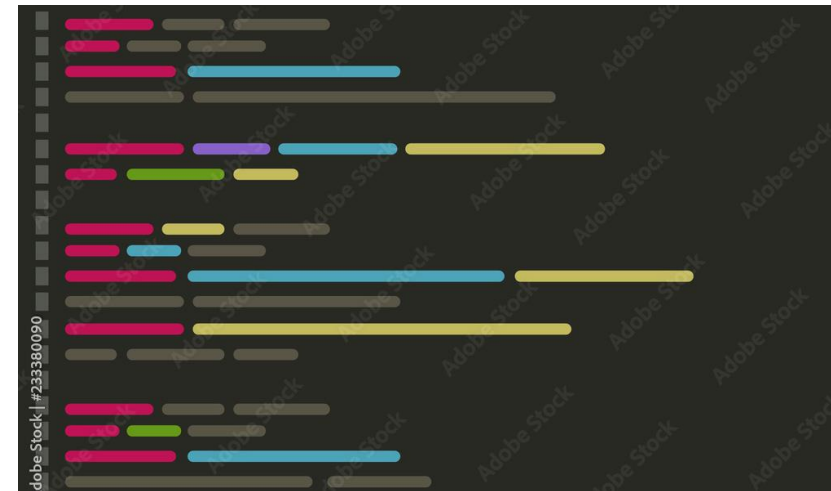

eclipse



- Unable to reproduce user's failures
- Configurability & Portability
- Difficult to foresee & reproduce

Why modern methods are ineffective

- Collect stack traces & register dumps.
- Limited for reproducing failure
- Identify correlations among crashes
- Possible solution: Collect only limited amounts of information like **Sample branch profiles**



Why modern methods are ineffective



- Usefulness \uparrow when details \uparrow
- Novel techniques (richer data)
- Limitations:
 - Expensive to collect
 - Bound to contain sensitive data

Example of the problematic

- Fault at line 30
 - One character skipped & extra character added in the output string
 - Index of output grow larger than the index of the input by one for each \\
- Unlikely achieved by random testing to reach POF

```
1. char replaceescape(char e) {
2.     switch (e) {
3.         case 'n':
4.             return 10;
5.         case 't':
6.             return 9;
7.         default:
8.             return 0;
9.     }
10. }
11. char uppercase(char l) {
12.     return l-'a'+'A';
13. }
14. void printresult(char* str,int length) {
15.     int i;
16.     for (i=0;i<length;i++) {
17.         if (str[i]!=0)
18.             printf("%c",str[i]);
19.     }
20. }
21. void process(char* source, char* dest) {
22.     int out=0;
23.     int in=0;
24.     int srclength = strlen(source);
25.     while (in<srclength) {
26.         if (source[in]>='a'&&source[in]<='z') {
27.             dest[out]=uppercase(source[in]);
28.         } else if (source[in]=='\\') {
29.             dest[out]=replaceescape(source[in+1]);
30.             out++; // correct version: in++;
31.         } else
32.             dest[out]=source[in];
33.         out++;
34.         in++;
35.     }
36.     printresult(dest,out);
37. }
38. int main(int argc, char *argv[]) {
39.     if (argc!=2)
40.         exit(0);
41.     if (strlen(argv[1])>=256)
42.         exit(0);
43.     char* outputstr=malloc(256);
44.     process(argv[1],outputstr);
45. }
```

Replace character by a new line or tab if \n or \t or null.

Returns the uppercase

Prints the result

If alphabetical and lower case turn it in uppercase

If it's escape run replaceescape

Figure 2. Example of faulty program.




The problematic



Goal of the study

Our goal

- Address the limitations by developing **novel approaches**
- E' should result in failure F'
- E' should execute P to produce F'
- Should generate E' only using P and D
- D should not contain sensitive data



What is BugRedux?

"At any point of a symbolic execution, called a symbolic state it is expressed as a path condition".

- A tool that implements a technique that
 - Collect execution data
 - Use the data to synthesize in-house executions
- **Symbolic Executions:** Executes a program using symbolic instead of concrete inputs.
- **Symbolic State:** Representation of a moment in a symbolic execution
- **Path condition (PC):** Conditions on the inputs that cause the execution to reach that point.

gatech/bugredux

A test repo for bugredux



1

Contributor

0

Issues

1

Star

2

Forks



What is BugRedux?

- **Symbolic execution** of path $\langle 1, 2, 3, 4, 6, 7, 8, 13 \rangle$
- **Symbolic state** at line 13 would be $\{[a \rightarrow a_0], [b \rightarrow b_0], [c \rightarrow c_0], [d \rightarrow a_0 + 4]\}$
- **Path condition (PC)** would be $(a_0 + 4 < b_0) \wedge (b_0 \leq 5) \wedge (a_0 < 5) \wedge (a_0 + 4 < c_0)$

```
function foo(int  $a_0$ , int  $b_0$ , int  $c_0$ ) {  
  ① int d = a + 4  
  ② if (d < b)  
  ③   //do something  
  ④ if (b > 5)  
  ⑤   //do something  
  ⑥ else if (a < 5)  
  ⑦   if (d < c)  
  ⑧   //do something  
  ⑨   else  
  10.   //do something  
  11. else  
  12.   //do something  
  ⑬ return  
}
```



The problematic



Goal of the
study



Methodology

The Specs

- 4 variations of BugRedux
 - Point of Failures (POF)
 - Stack traces
 - Call sequences
 - Complete program traces
- Implemented in C
- Applied on 17 failures of 15 real world programs



The architecture

- **Instrumenter**: Application → Instrumented application
- **Analyzer**: Crash reports → Test input
- **Software Tester**: Test input → Recreate field failure

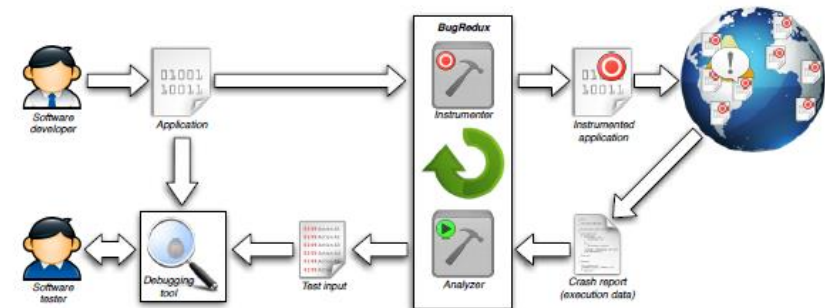


Figure 4. Intuitive high-level view of BUGREDUX.

The algorithms

- *GenerateInputs*: Symbolic execution to identify set of intermediate goals.
- *Sel/NextState*: Heuristic based on distance to select states

Algorithm 2: SelNextState

Input : *icfg* : ICFG for program *P*

states_set: set of symbolic states

curr_goal: next goal

Output: *ret_state*: candidate state for exploration

```

1 begin
2   mindis  $\leftarrow +\infty$ 
3   ret_state  $\leftarrow$  null
4   foreach Statei  $\in$  states_set do
5     if Statei.goal == curr_goal then
6       if Statei.loc can reach curr_goal in ICFG then
7         nd  $\leftarrow$  shortest distance from Statei.loc to
          curr_goal in ICFG
8         if nd < mindis then
9           mindis  $\leftarrow$  nd
10          ret_state  $\leftarrow$  Statei
11        end
12      end
13    end

```

Algorithm 1: GenerateInputs

Input : *icfg* : ICFG for program *P*

goals_list : an ordered list of statements *G₀*, ...*G_n*

Output: *input_f*: candidate input for synthesized run

```

1 begin
2   sym_state0  $\leftarrow$  initial symbolic values of program inputs
3   states_set  $\leftarrow$  (icfg.entry, true, sym_state0, G0)
4   curr_goal  $\leftarrow$  G0
5   while true do
6     curr_state  $\leftarrow$  null
7     while curr_state == null do
8       curr_state  $\leftarrow$ 
        SelNextState(icfg, states_set, curr_goal)
9       if curr_state == null then
10        if curr_goal  $\neq$  G0 then
11          curr_goal  $\leftarrow$  previous goal in goals_list
12          continue
13        else
14          return null
15        end
16      end
17    end
18    if curr_state.cl == curr_goal then
19      if curr_goal == Gn then
20        inputf  $\leftarrow$  solver.getSol(curr_state.pc)
21        if inputf is found then
22          return inputf
23        else
24          remove(curr_state, states_set)
25          continue
26        end
27      else
28        curr_goal  $\leftarrow$  next target in goals_list
29        curr_state.goal  $\leftarrow$  curr_goal
30      end
31    else
32      if curr_state.cl  $\in$  goal_list then
33        remove(curr_state, states_set)
34        continue
35      end
36    end
37    if curr_state.cl is a conditional statement then
38      curr_state.pc  $\leftarrow$ 
        addConstr(curr_state.pc, pred, true)
39      curr_state.cl  $\leftarrow$  getSucc(curr_state.cl, true)
40      if solver.checkSat(curr_state.pc) == false then
41        remove(curr_state, states_set)
42      end
43      false_pc  $\leftarrow$ 
        addConstr(curr_state.pc, pred, false)
44      false_cl  $\leftarrow$  getSucc(curr_state.cl, false)
45      if solver.checkSat(curr_state.pc) == true then
46        new_state  $\leftarrow$ 
          (false_cl, false_pc, curr_state.ss, curr_state.goal)
47        insert(new_state, states_set)
48      end
49    else
50      curr_state.ss  $\leftarrow$ 
        symEval(curr_state.ss, curr_state.cl)
51      curr_state.cl  $\leftarrow$  getSucc(curr_state.cl)
52    end
53  end
54 end

```




The Research Questions

RQ1: Can BugRedux synthesize executions that are able to reproduce field failures starting from a set of execution data?

RQ2: If so, which types of field data provide the best tradeoffs in terms of cost benefit?

The Data

- 4 variations of BugRedux
 - POFs (no modif)
 - Call Stacks (no modif)
 - Call Sequences (instruments call sites)
 - Complete program traces (instruments all branches)
- POFs: Single entry
- Call Stacks: Entry for each function
- Call Sequences: *goals_list* would contain entry for each call
- Complete traces: *goal_list* would be an entry per branch



Methodology

- 3 public repositories:
 - 3 from SIR
 - 2 from BugBench
 - 9 from exploit-db
- Chose because of past researches and representativeness
 1. Ran test cases until fail
 2. Run all the failing 3 times
 - 2.1 Ran on unmodified program
 - 2.2 Ran on BugRedux to collect Call Sequences
 - 2.3 Ran on BugRedux to collect complete traces
- Collect time and size of execution data generated
- Use 4 variants of BugRedux to synthesize candidate input
- See if inputs can reproduce the original failure
- Timeout of 72 hours for generation

Table I
SUBJECT PROGRAMS USED IN OUR STUDY.

Name	Repository	Description	Size (kLOC)	# Faults
sed	SIR	stream editor	14	2
grep	SIR	pattern-matching utility	10	1
gzip	SIR	compression utility	5	2
ncompress	BugBench	(de)compression utility	2	1
polymorph	BugBench	file system “unixier”	1	1
aeon	exploit-db (CVE-2005-1019)	mail relay agent	3	1
iwconfig	exploit-db (CVE-2003-0947)	wireless tool	11	1
glftpd	exploit-db (OSVDB-ID#16373)	FTP server	6	1
htget	exploit-db (CVE-2004-0852)	file grabber	3	1
socat	exploit-db (CVE-2004-1484)	multipurpose relay	35	1
tipxd	exploit-db (OSVDB-ID#12346)	IPX tunneling daemon	7	1
aspell	exploit-db (CVE-2004-0548)	spell checker	0.5	1
exim	exploit-db (EDB-ID#796)	message transfer agent	241	1
rsync	exploit-db (CVE-2004-2093)	file synchronizer	67	1
xmail	exploit-db (CVE-2005-2943)	email server	1	1



The problematic



Goal of the study



Methodology



Results

Results

Table II Overhead by BugRedux

TIME OVERHEAD IMPOSED ON THE SUBJECT PROGRAMS.

Name	POF	Call stack	Call sequence	Complete trace
sed.fault1	0	0	4.5%	27.2%
sed.fault2	0	0	12.5%	87.5%
grep	0	0	47%	182%
gzip.fault1	0	0	10.3%	72%
gzip.fault2	0	0	12%	308%
ncompress	0	0	2%	16%
polymorph	0	0	1%	8%
aeon	0	0	50%	1066%
iwconfig	0	0	7%	128%
glftpd	0	0	9%	45%
htget	0	0	9%	287%
socat	0	0	21%	110%
tipxd	0	0	2%	36%
aspell	0	0	18.8%	143%
exim	0	0	17.4%	389%
rsync	0	0	3%	66%
xmail	0	0	22.6%	290%

Depends on program size and execution length

Table III Function call low but loops high

SIZE OF EXECUTION-DATA COLLECTED (KB).

Name	POF	Call stack	Call sequence	Complete trace
sed.fault1	0.8	0.8	5.8	54.4
sed.fault2	0.9	0.9	10.2	261.9
grep	0.7	0.7	3.4	716.1
gzip.fault1	0.8	0.8	2.0	176
gzip.fault2	0.8	0.8	2.5	1784.6
ncompress	0.7	0.7	0.9	33.1
polymorph	0.5	0.5	0.7	1.5
aeon	1	1	1.1	3
iwconfig	1	1	3	594
glftpd	1.5	1.5	3.2	130
htget	0.7	0.7	2.7	2814
socat	0.8	0.8	9.6	451
tipxd	0.6	0.6	0.7	19
aspell	0.6	0.6	30.5	566
exim	0.9	0.9	100.7	14897
rsync	1	1	11.4	521
xmail	0.8	0.8	84.8	2361

Same (Extract from crash reports together)²⁰

Results

Table IV

Leaves more degree of freedom
& ↑ chance of success

EFFECTIVENESS AND EFFICIENCY OF BUGREDUX IN SYNTHESIZING
EXECUTIONS STARTING FROM COLLECTED EXECUTION DATA.

Name	POF	Call stack	Call sequence	Whole trace
sed.fault1	N/A	N/A	98s	N/A
sed.fault2	N/A	N/A	17349s	N/A
grep	N/A	16s	48s	N/A
gzip.fault1	3s	18s	11s	N/A
gzip.fault2	20s	28s	25s	N/A
ncompress	155s	158s	158s	N/A
polymorph	65s	66s	66s	N/A
aeon	1s	1s	1s	1s
rsync	N/A	N/A	88s	N/A
iwconfig	5s	5s	5s	N/A
glftpd	5s	5s	4s	N/A
htget	53s	53s	9s	N/A
socat	N/A	N/A	876s	N/A
tipxd	27s	27s	5s	N/A
aspell	5s	5s	12s	N/A
xmail	N/A	N/A	154s	N/A
exim	N/A	N/A	N/A	N/A

Provide more
information

Table V

EFFECTIVENESS OF SYNTHETIC EXECUTIONS IN REPRODUCING
OBSERVED FAILURES.

Name	POF	Call stack	Call sequence	Complete trace
sed.fault1	N	N	Y	N
sed.fault2	N	N	Y	N
grep	N	N	Y	N
gzip.fault1	Y	Y	Y	N
gzip.fault2	N	N	Y	N
ncompress	Y	Y	Y	N
polymorph	Y	Y	Y	N
aeon	Y	Y	Y	Y
rsync	N	N	Y	N
iwconfig	Y	Y	Y	N
glftpd	Y	Y	Y	N
htget	N	N	Y	N
socat	N	N	Y	N
tipxd	Y	Y	Y	N
aspell	N	N	Y	N
xmail	N	N	Y	N
exim	N	N	N	N



- 1) POF are close to entry
- 2) Failures triggered by reaching POFs

Now for the questions...

- **RQ1: Can BugRedux synthesize executions that are able to reproduce field failures starting from a set of execution data?**
- Yes
- **RQ2: If so, which types of field data provide the best tradeoffs in terms of cost benefit?**
- Call Sequences (acceptable time & space overhead and hides confidential information about an execution)
- Downside of POF and Call stack: POF distant from the fault, which makes POF and Call stack irrelevant.
- Downside of Complete Trace: Providing too much guidance may be just as problematic as providing too little. (Sweet spot)

What we learned

Richest data, beside being expensive to collect, is not necessarily the most useful when used for synthesizing executions.

Best cost-benefit ratio is
Call Sequences
(recreates 16 out of 17 failures)



To conclude

- Recreating failures in the field is arduous
- To address this problem, we presented BugRedux
 - Collecting data about program runs in the field
 - Extract the data collected to mimic the failing execution & reproduce failure
- We found that using Call Sequences is the most efficient way to reproduce failures